

Definition

Project Overview:

Although recognizing characters in documents can be considered a solved task for computer vision [1], the same task is still considered challenging in the context of natural scenes: characters and digits are usually corrupted by several phenomena like blur, distortion and illumination effects among others [2]. Manually detecting and transcribing character string in billions of images, however, is a daunting task; having the capability to correctly and automatically recognize characters directly from raw images would enable and encourage a wide range of real world applications such as mapping and geocoding [3], captcha recognition [4] and license plate recognition [5].

Netzer et al. [2] introduced a digit classification dataset SVHN¹ captured from Google Street View images which includes over 600000 digit samples. They also proposed a recognition approach where individual characters would be segmented out of the original image and a post-classification model would be employed to recognize these individual characters. More recently, however, Goodfellow et al. [3], introduced a deep convolutional architecture capable of unifying this procedure while operating directly on the image pixels.

In this project, I have created a Python application capable of recognizing number strings from real-world images taken from a webcam. The application uses a Deep Convolutional Network trained using the SVHN dataset and following the architecture and approach described by Goodfellow et al. [3].

Problem Statement:

The goal of the project is to develop an application capable of identifying sequence of numbers within real-world images without the need of segmenting single digits beforehand. I would like to avoid the use of sliding window approaches, while also

¹ The Street View House Numbers (SVHN) Dataset. <http://ufldl.stanford.edu/housenumbers/>

keeping the network complexity low (due to computational power limitations). Because of this, I will implement a 2-stage recognition system as the one shown in Figure 1.




Figure 1 Implemented 2-stage recognition model

The job of the regression model will be to identify the location of the numbers within the image. It will output a set of coordinates corresponding to the bounding box containing the number string. I will crop this bounding box and I will pass it to a second stage (Classification Model), which will take care of recognizing which numbers can be seen in the image. I will train models which should be capable of recognizing number strings composed of up to 5 digits.

I will approach this problem by following these steps:

1. Generate a synthetic dataset by concatenating character images from the MNIST² database of handwritten digits.
2. Design and train a network architecture capable of achieving good performance on the synthetic dataset.
3. Create a more challenging and realistic scenario by downloading and preprocessing the SVHN dataset.
4. Train the model with this new dataset and evaluate its performance.
5. Design and train a regression model using the bounding boxes provided by the SVHN dataset.
6. Develop an interface to load video frames from a web camera and to identify number strings within those frames.

Metrics:

Like Goodfellow et al. [3], I will determine the *accuracy* of our Classification Model by computing the proportion of the testing images for which the length of the sequence and every element or digit in the sequence is predicted correctly. There

² The MNIST database of handwritten digits. <http://yann.lecun.com/exdb/mnist/>

will be no “partial credit” for correctly classifying individual digits. Using this metric will allow me to compare the performance of the system to the one proposed by Goodfellow et al. [3].

The performance of our Regression Model, however, will be measured in terms of the *Dice's coefficient* Q_s , which measures the similarity between the predicted Bounding Box (X) and the ground truth (Y). $|X \cap Y|$ represents the overlapping area between the predicted and the real bounding boxes. A coefficient of 1 represents a perfect match.

$$Q_s = \frac{2|X \cap Y|}{|X| + |Y|}$$

I will also compute the *accuracy* with which the Regression Model is capable of identifying whether or not there is a number in the image.

Analysis

Data Exploration:

1. THE MNIST DATABASE OF HANDWRITTEN DIGITS

The MNIST database provides over 60000 examples of individual handwritten digits. The digits have been size-normalized and centered in a fix-sized image (28x28 pixels). Figure 2 shows some examples of the MNIST dataset images. In the next section, I will describe how did I use this dataset to generate synthetic number strings for training my first classification model.

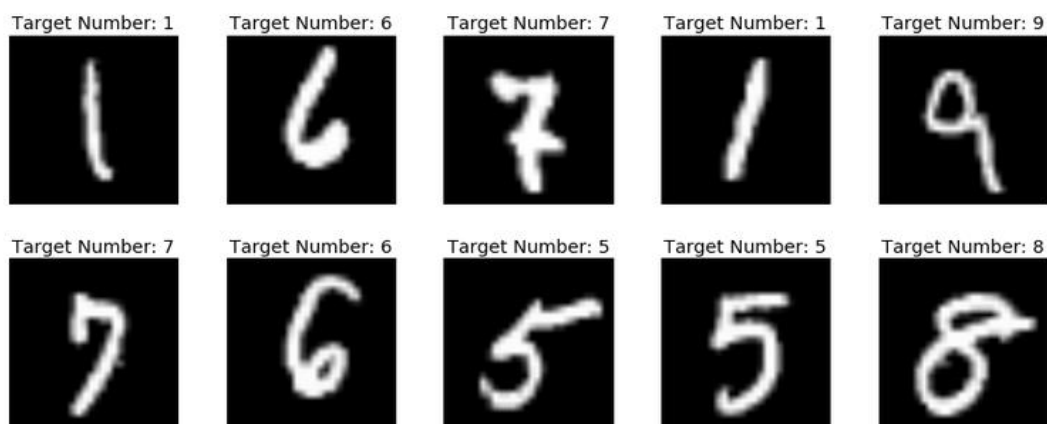


Figure 2 Some examples from the MNIST dataset.

2. THE STREET VIEW HOUSE NUMBERS (SVHN) DATASET

The SVHN dataset was obtained from house numbers in Google street view images and it represents a more real and challenging benchmark when compared to the MNIST database. Figure 3 shows the type of images which are included in this dataset.



Figure 3 Images from the SVHN dataset. Taken from <http://ufldl.stanford.edu/housenumbers>

The bounding boxes in Figure 3 are shown just for illustration purposes given that this information is stored in an additional file. The downloaded folders contain both the original images in *png* format and a *digitStruct.mat* file. The *digitStruct.mat* file contains a structure called ***digitStruct*** which is organized as follows:

- **digitStruct** [struct]: structure with the same length as the number of images
 - **name** [string]: filename of the corresponding image
 - **bbox** [struct]: structure containing the bounding boxes and labels
 - **label** [int]: number in the current bounding box
 - **height** [int]: height of the bounding box
 - **width** [int]: width of the bounding box
 - **left** [int]: x position of the top left corner
 - **top** [int]: y position of the top left corner

Exploratory Visualization:

1. GENERATING A SYNTHETIC DATASET

As suggested in the Deep Learning Capstone Project, I decided to generate a synthetic dataset by concatenating single digits taken from the MNIST database. The code randomly selects images from the MNIST dataset and concatenates them as shown in Figure 4. To generate different types of images, the script also randomly selects a *shifting factor* δ , which determines the slope of the number string.



Figure 4 Synthetic dataset generation process. N digits and one shifting factor δ are randomly selected. The images are then concatenated after shifting the location of the second digit according to δ . Empty spaces are filled with black. The resulting image is then resized to 32x32 pixels.

The maximum length of the number strings (N) can be easily modified on code. For testing purposes, I generated numbers with up to $N=5$ digits. The script was designed to generate a homogeneous dataset containing a similar number of images with different string lengths as shown in Figure 5.




Figure 5 Distribution of the number string lengths in the three different synthetic datasets

Figure 6 shows some of the synthetic images which were generated using the code described above. All the images were saved in grayscale after subtracting their mean.




Figure 6 Synthetic images generated by concatenating single digits from MNIST database

The number of synthetic images which are generated can be also modified on code. In total, 200000 images were generated this way from which 80% were assigned to training, 10% to validation and 10% for testing.

2. THE STREET VIEW HOUSE NUMBERS (SVHN) DATASET

The SVHN dataset is composed not only by training and testing samples but also by an additional, somewhat less difficult, *extra set* which can be used as additional training. The number of images which are available on each dataset is shown below.

Folder	Number of images
Training	33402
Testing	13068
Extra	202353

Because the dataset is composed by variable-resolution images, it is interesting to get some statistics regarding the size of the available images:

	Training	Testing	Extra
Avg. Height	57,21	71,56	60,80
Max. Height	501	516	415
Min. Height	12	13	13
Avg. Width	128,28	172,58	100,38
Max. Width	876	1083	668
Min. Width	25	31	22

As it can be seen in the table above, there is a wide range of image sizes in the available datasets. This is good because it will provide us with a lot of resolution and scaling variability.

It is also interesting to analyze the length of the number strings which can be found in the available images. If all of them contained only one digit, it would not be possible to train a multi-digit classifier. Knowing the distribution of number string lengths can help us to determine if the dataset is useful for us or not. **Error! Reference source not found.** shows the percentage of the datasets which contains number strings with a given length.




Figure 7 Distribution of the number string lengths in the three different SVHN datasets

From Figure 7 we can determine that the datasets are composed mainly by number strings with up to 4 digits. The number of images with 5 digits, a total of 126 samples, is negligible and there is only one image presenting more than 5 digits.

Based on these results I will design the network to deal with up to 5 digits. Given the lack of training examples, however, I do not expect the network to perform well on images containing 5 numbers.

Algorithms and Techniques:

I will try to replicate the classifier presented by Goodfellow et al. [3] which is based on a Convolutional Neural Network (CNN). CNNs can represent very complex models and are currently the state-of-the-art technique for most image processing tasks. CNNs are very similar to regular neural networks but their architecture is designed to take advantage of the 2D structure of the input image (or a 2D input such as audio signals). A typical CNN architecture consists of several convolutional and subsampling layers followed by fully connected layers.

The reason why convolutional layers have become the most employed technique in computer vision applications is because of their capability of *sharing* parameters. It is impractical, when dealing with high dimensional inputs, to connect new layers to all the neurons in the previous volume. As shown in Figure 8, convolutional layers connect each neuron to only a local region of the input volume, thus creating a set of learnable filters. This is possible under the reasonable assumption that if one feature is useful at some spatial position (x_1, y_1) , then it should also be useful at another location (x_2, y_2) . During operation, each filter is convolved across the height and width of the input volume, thus obtaining 2D activation maps which represent the response of each one of these filters at every spatial location. During training, the network tunes each one of these filters according to the task at hand, thus autonomously finding a set of relevant and useful features and making hand-crafted features obsolete.




Figure 8 An input volume in red (e.g. 32x32x3 pixel image) and the volume of neurons in one convolutional layer (blue). Each neuron in the convolutional layer (5 in this example) is connected only to a small region in the original input, but to the full depth. Taken from <http://cs231n.github.io/convolutional-networks/>.

There are several parameters which need to be tuned when optimizing this kind of architectures. Although some of these parameters will be tuned with the help of a *Validation Set*, many others will be chosen based on state-of-the-art architectures, recent literature and computational constraints.

There are different ways of approaching the current recognition problem and a 2-stage solution as the one presented in Figure 1 is not strictly necessary. I could train a complex model capable of coping with scale and location variances, but that would require a lot of data, a lot of parameters and therefore, a high computational power and training time. A second approach could involve a sliding window technique, but I decided to go for something a bit more interesting and use a second convolutional network instead.

The first network (Regression Model) will take the original image (properly scaled) as input and it will output the following vector:

Index	Meaning
Output [0:1]	Probability of having a number string within the provided image (confidence score)
Output [2:5]	Bounding Box surrounding the number [Top, Left, Bottom, Right]

If a number is detected within the image, the system will crop the corresponding bounding box. This crop will be properly resized and passed to the second network (Classification Model), which in turn will output a classification probability for up to 5 digits. The Classification Model will consider 11 different classes: 0-9 to represent the corresponding digits, and 10 to represent empty spaces.

Benchmark:

Goodfellow et al. [3] achieved a sequence transcription accuracy of 96.03% on the SVHN dataset after training for about 6 days(!). My goal, however, will be to implement some more recent techniques to reduce the model complexity without affecting its performance too much. Given that I cannot compete with the kind of computational power which is usually employed when training deep models, I will focus on optimizing the network architecture. My ideal performance would be obtaining a transcription accuracy above 90%. As it was mentioned in the previous sections, there will be no “partial credit” for correctly classifying individual digits and a transcription will only be considered correct if all the digits are classified correctly.

Methodology

Data Preprocessing:

1. PREPROCESSING SVHN DATASET (CLASSIFICATION MODEL)

I performed a very similar preprocessing as the one described by Goodfellow et al. [3]:

- I first found the smallest rectangular bounding box containing all the digits in the respective image. This can be easily done by analyzing the individual bounding boxes provided by the SVHN dataset shown in red to the left of Figure 9.

- I then expanded this bounding box by 30% in both x and y directions, crop the image to that bounding box and resized the crop to 64x64 pixels. One of this expanded bounding boxes is shown in blue to the left of Figure 9.
- After that, I cropped several 54x54 pixel images from random locations within the previous 64x64 image. Figure 9 (right) shows some of the 54x54 crops made over the 64x64 pixels image. Although Goodfellow et al. [3] reported an improvement of only 0.5% using this data augmentation, I believe this to be an important step to achieve good performance with real images.



Figure 9 To the left, the individual bounding boxes provided by the SVHN dataset are shown in red together with the expanded bounding box in blue. To the right, the image has been cropped and resized, and 54x54 pixel images are cropped from random locations (shown in red and blue)

All the images are then converted into grayscale to reduce the model complexity. Finally, we subtract the mean of each image and apply random brightness and contrast transformations.

2. PREPROCESSING SVHN DATASET (REGRESSION MODEL)

I used the SVHN dataset once again to generate training and validation samples for the Regression Model. To guarantee scale and location variance in our training samples, I generated several images by sliding a window through the original images. The stride of the sliding window can be easily modified on code. Figure 10 shows 3 training samples generated from the original image. More than 300000 training images were generated this way.



Figure 10 Examples of data augmentation for the Regression Model. Generating several samples by shifting the location of the bounding boxes and changing their scales.

Because our Regression Model also outputs the probability of the image containing a number, we required some “negative” training examples. I employed the CIFAR-10³ dataset which contains 50000 images of 10 different categories like birds, cats, airplanes and dogs among others. The final dataset was split in training, validation and testing sets.

The Regression Model, in contrast to the Classification Model, employs RGB images and linearly scales them to have zero mean and unit norm (image whitening).

Implementation:

All the code was implemented as Jupyter Notebooks. The code, both for the Regression and the Classification models, was divided into the following steps:

1. Both the validation and test sets are loaded into memory. The preprocessing described before is applied to all samples.
2. The training dataset, on the contrary, had to be read dynamically from binary files due to computational constraints. Loading the data this way allowed me to spare a lot of memory during training.
3. Define accuracy function.
4. Define the network architecture.
5. Define loss function and learning method.
6. Train the network logging training and validation loss into Tensorboard.
7. Save the model.
8. Plot the performance on some sample images.

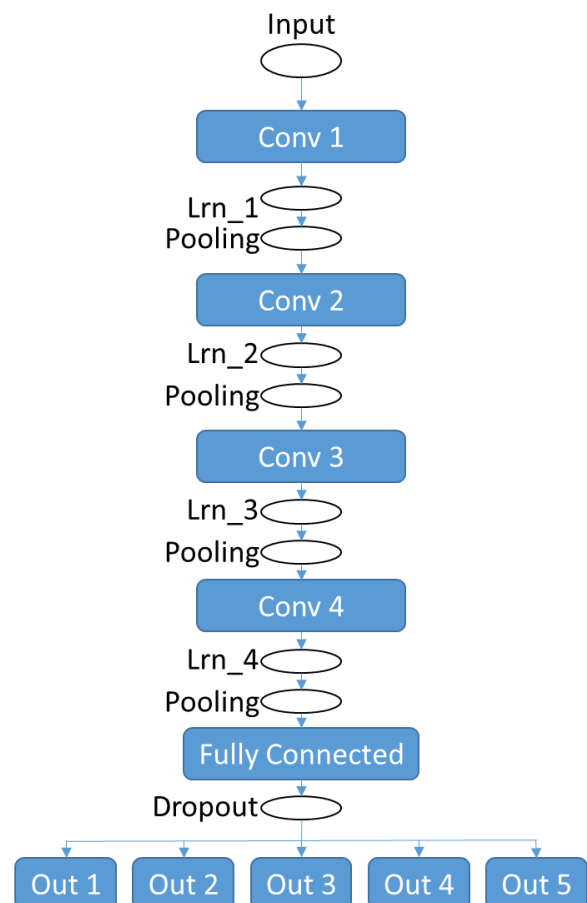


Figure 11 Networks architecture.

³ The CIFAR-10 dataset. <https://www.cs.toronto.edu/~kriz/cifar.html>

Both the Regression and the Classification networks share the same architecture shown in Figure 11 (although the Regression Model employs an additional output). The specific details of these networks will be presented in the following sections.

The loss of the Classification Model is defined as the mean cross entropy error between the outputs and the labels. The loss of the Regression Model is defined as the mean L2 distance between the bounding box coordinates plus the mean cross entropy error of the confidence score.

Both networks employ an exponentially decaying learning rate. The initial learning rate was tuned for each network individually.

Regarding the learning method, I employed an *Adagrad Optimizer* [6] to minimize the loss on both networks. As it will be shown in the following section, I choose *Adagrad* after comparing its performance with a couple of different optimization methods. *Adagrad* is a gradient-based optimization algorithm widely used to train deep networks. Its main characteristic is that it adapts the learning rate to the parameters, thus performing larger updates for infrequent and smaller updates for frequent parameters [7]. It has been found that *Adagrad* greatly improve the performance of *stochastic gradient descent* (SGD). One of *Adagrad*'s main benefits is that it eliminates the need to manually tune the learning rate.

The training steps can be modified on code and all the parameters are saved periodically during, and at the end of the training. It is possible to run the training code several times to improve the performance of current models.

Once both networks are successfully trained, it is possible to test the system on real webcam images. I have provided a final Jupyter Notebook which uses OpenCV to load camera frames as a video sequence. The original images are resized to 48x64 pixels (I choose these values because my webcam provides 480x640 images and it was therefore possible to avoid distortions) and they are used as input to the Regression Model. If this first network detects the presence of a number within the image, the system crops the image according to the predicted bounding box. This crop is then sent to the second network which converts it to grayscale and which recognize the numbers in the image.

The provided code displays the captured images together with the predicted bounding box (which is scaled to the original image size) and the number which is detected on the top left corner as shown in Figure 12.



Figure 12 Evaluating the networks performance directly on a webcam frame. The predicted bounding box is shown in blue and the number transcription is drawn at the top left corner.

Refinement:

I started using the Classification architecture shown to the right. All the convolutional layers have a stride of one and are denoted as:

Conv{receptive field size}-{number of channels}

The ReLU activation functions are not shown for brevity. Additionally:

- Max-pooling is performed over a 2x2 pixel window, with stride = 2.
- The batch size was set to 64 samples.
- The dropout probability was set to 0.95 during training.
- The initial learning rate α_0 was set to 0.04 with a decaying factor of 0.95.

Initial Configuration
Input (32x32 Grayscale Model)
Conv5-16
Local Response Normalization
Max Pooling
Conv5-32
Local Response Normalization
Max Pooling
Conv5-64
FC-64
Dropout
5 x Soft-Max

Figure 13 shows the performance of this initial model when trained and tested using the generated synthetic dataset in comparison to the performance obtained when training and testing the model with the real, previously preprocessed, SVHN dataset. The initial architecture performs quite well when using the synthetic dataset achieving a validation accuracy of 89.7% after 20000 training steps. The same model, however, performs poorly on the more demanding SVHN dataset, achieving a validation accuracy of 82.9%.




Figure 13 Performance of the initial model when trained and tested using the synthetic or SVHN datasets.

In contrast, the model proposed by Goodfellow et al. [3] consisted of eight convolutional layers and significantly more parameters, obtaining a transcription accuracy of 96.03%. Unfortunately, I do not have the computational resources to train such a model.

For this reason, I decided to implement some of the ideas proposed by Simonyan & Zisserman [8] in their VGG architecture. The main idea is that we can replace a large receptive field convolutional layer with a stack of very small convolutional filters. They demonstrated that doing this not only reduces the number of parameters in the model, but also improves the performance by making the decision function more discriminative. As I will show in the following section, I replaced all the 5x5 conv. layers by a stack of two 3x3 layers.

This was, however, not the only refinement I had to make to improve the performance of the network. As mentioned before, I decided to use an *Adagrad* optimizer after comparing its performance with other two techniques: *Momentum* and *RMSprop*. **Error! Reference source not found.** shows the performance of the initial model on the synthetic validation set when trained using the three different optimization methods. The initial learning rate was set to 0.04 during the three experiments. The momentum term γ was set to 0.9. Any additional tuning parameter was left as the default values provided by TensorFlow. This is clearly not an exhaustive evaluation, but it allowed me to relatively quickly compare the results and determine which optimizer would work better with my model without the need of intensive tuning. As it can be seen in **Error! Reference source not**

found., *Adagrad* presents the best performance without the need of additional tuning. *RMSprop*, on the contrary, performs quite poorly when implemented with the default values, which makes me think that additional tuning would be required if I were to choose this method.



Figure 14 Comparing different optimization techniques with respect to the validation accuracy when training on the synthetic dataset.

The following parameters were also modified in an iterative way trying to optimize the performance on the validation set while avoiding overfitting:

- Number of convolutional layers: I tried models with up to 5 convolutional layers (each one consisting of a stack of 2). Having 5 layers did not improve the performance significantly, but it did the training process much slower. The final model employs 4x2 layers.
- Depth of the convolutional layers and nodes in the fully connected layer: I increased the number of channels on each layer and the size of the FC layer slowly trying to find a good threshold between performance and training time.
- Learning rate: I had some problems when starting the project and it was all due to a large learning rate. The model was overshooting during training and the loss was not decreasing even after training for many steps. It took me some time to find out the reason, but once I did I could find optimal values.
- Dropout: The dropout probability was also tuned after trying a couple of different values.

Results

Model Evaluation and Validation:

As it has been mentioned before, I used validation sets to evaluate the performance of both models and to properly configure their architectures. Figures 15 and 16 show a plot of the training and validation losses for both models. The training process was stopped whenever no further improvement was observed in the validation loss/accuracy. For this reason, there does not seem to be any signs of overfitting in neither of both models.

Although it looks like the Regression model stops learning at around 30000 iterations (Figure 16), additional training steps helped improving the Dice's coefficient which the model achieves in both training and validation sets. Another thing to notice is that the loss of both models does not seem to decrease during the first ~6K steps. This could be happening due to a large initial learning rate and additional tuning could still improve the training process.

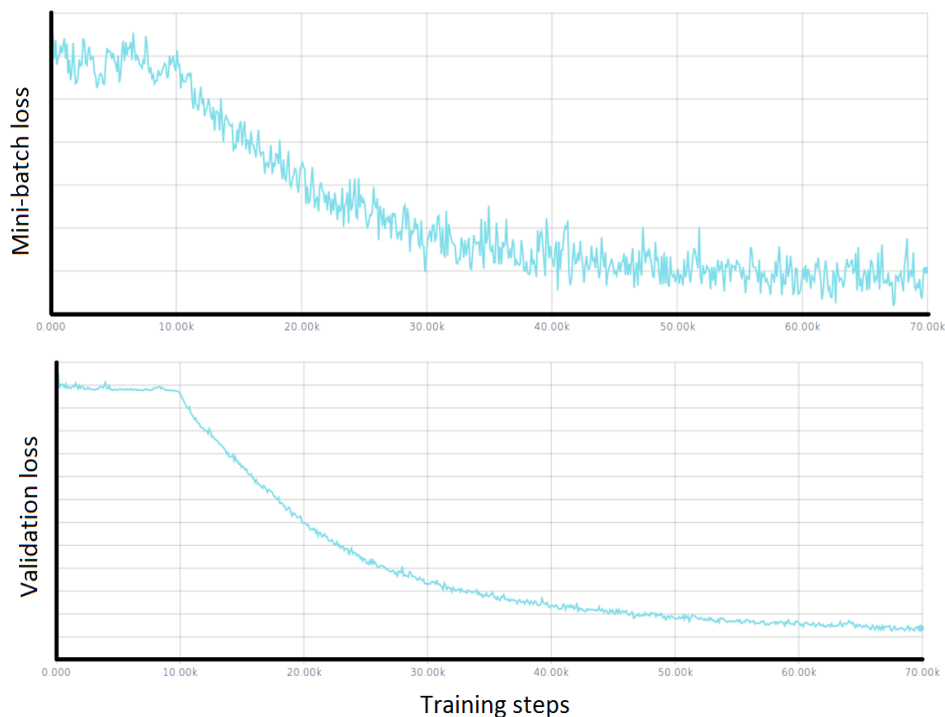


Figure 15 Training and validation loss presented by the Classification model. Training was stopped when no further improvement in the validation set was observed.




Figure 16 Training and validation loss presented by the Regression model.

The final configurations are presented in the following table where I use, once again, the notation: “Conv{receptive field size}-{number of channels}”. I trained different models and the ones presented here were the ones which performed the best.

ConvNets Configuration	
Regression Model	Classification Model
Input (48x64 RGB Image)	Input (54x54 Grayscale Image)
Conv3-32	
Conv3-32	
Local Response Normalization	
Max Pooling	
Conv3-48	
Conv3-48	
Local Response Normalization	
Max Pooling	
Conv3-64	
Conv3-64	
Local Response Normalization	
Max Pooling	
Conv3-80	Conv3-96
Conv3-80	Conv3-96
FC-2048	FC-3136
Dropout	
[Soft-Max] + [4 x ReLU]	5 x Soft-Max

All the convolutional layers have a stride of one, so that neither of them decreases the resolution of the image. The ReLU activation functions are not shown for brevity. Additionally:

- Max-pooling is performed over a 2×2 pixel window, with stride = 2.
- The batch size was set to 64 samples.
- The dropout probability was set to 0.9375 during training.
- The initial learning rates α_0 were set to 0.002 and 0.015 for the Classification and Regression models respectively, with a decaying factor of 0.95.

The final performance of both models is presented in the next table. The Classification network achieves a transcription accuracy of almost 90% on the testing dataset. Similarly, the Regression network reaches a detection accuracy of about 84% while achieving an average Dice's coefficient of 0.62. I noticed, however, that the SVHN testing samples are considerably more complex when compared to the training ones. For this reason, I also evaluated the network on a separate testing set which was generated using some samples taken from the SVHN training folder; images which of course were removed from training datasets. The Regression model reaches a Dice's coefficient of 0.81 on this additional testing set.

Evaluation Dataset	Classification model	Regression model	
	<i>Transcription accuracy (%)</i>	<i>Detection accuracy (%)</i>	<i>Dice's coefficient</i>
Training	90	100	0.87
Validation	89.8	85.5	0.62 (*0.82)
Testing	89.2	83.9	0.62 (*0.81)

Justification:

The final architecture achieves a transcription accuracy of 89.2% which is very close to the initially established benchmark (90%). Considering the difference in complexity, number of parameters and training time required by our models in comparison to the one proposed in [3], I am very satisfied with this performance.

Moreover, as presented in Figure 17, the robustness of the proposed approach was evaluated using real webcam images. The system localizes and classifies the numbers which are captured by the camera and it is very robust to changes in orientation and scale. Besides, the system runs online, thus making transcriptions directly from video sequences.

The system, however, is far from perfect. There are some cases where the networks fail either to localize the numbers or to correctly classify them. Some of these examples are shown in Figure 18. I noticed that shaking the objects I was holding during the video recordings had a strong and negative effect on the performance. Holding the numbers steadily and making slow movements resulted in a much better localization/classification accuracy.

It is important to note that although I have employed a relatively simple network configuration (in comparison to state of the art architectures), the performance of the system is quite decent. There are, of course, many ways to further improve the models as it will be discussed in the Improvement section.

Conclusion

Free-Form Visualization:

Figure 17 and 18 show some examples of correctly and incorrectly classified images respectively. The first thing to notice is that the Regression model operates quite well when the numbers are large, but it may fail as the numbers become smaller. This could be improved by further augmenting the training datasets by including samples where the number strings are relatively small. The Regression model is also quite robust to the location of the numbers within the image, which can be attributed to the way the training samples were generated. As shown in Figure 18, however, one clear problem in the Regression model is the presence of False Positives. Training the network with many more “negative” examples may help us to alleviate this problem.

The Classification model does also perform very well when the numbers are correctly localized by the Regression network. Nevertheless, I noticed that the network usually mistakes the number 3 for the number 8 (and vice versa). In other cases, the network may misclassify some numbers due to the presence of background patterns (like the vertical white lane to the left of the number 45, which can be mistaken for the number 1).




Figure 17 Examples of successfully classified images. The Regression models correctly identifies the location of the numbers and the Classification model is then capable of recognizing the numbers inside the bounding box.




Figure 18 Some examples where the networks fail either to localize or classify the numbers. In the first example, starting from the left, the network correctly localizes the numbers but it does not classify all the digits correctly. On the second example, the network cannot detect the number. The third image shows a false positive and, in the last example, the network incorrectly localizes the number string.

Reflection:

I have presented a 2-stage approach for detecting, localizing and classifying numbers from image inputs. I used a first Regression network to detect and localize the bounding box surrounding the numbers within the image, and a second Classification network to recognize which digits can be observed inside this bounding box. I have provided some algorithms to preprocess publicly available image datasets like MNIST and SVHN, and I have made use of data augmentation techniques to generate the additional training samples required by the proposed models.

The main reason why I choose to work on the Deep Learning Capstone Project was because of my desire to acquire practical experience implementing deep models and to get the opportunity to use and master libraries like TensorFlow. Because of this, I am very satisfied with the things I have learned while developing this application and very excited to start new and more challenging projects.

Improvement:

It is clear that the proposed solution could be significantly improved by increasing the complexity of the network architectures (make the networks deeper!). This will,

however, require many more training samples (to avoid overfitting), computational power and training time.

Nevertheless, there are some things which could be done without drastically changing the current system. I recently read that the pooling layers may decrease the performance of our Regression model, given that most of the spatial information is thrown away. We could redesign the Regression architecture by removing the pooling layers and properly adjusting the convolutional ones.

I also believe that some of the false positives delivered by the current systems are caused by the lack of sufficient “negative” samples during training. Once again, increasing our training dataset could be beneficial.

References

- [1] Pierre Sermanet, Soumith Chintala, TannLeCun. “Convolutional Neural Networks Applied to House Numbers Digit Classification”. In International Conference on Pattern Recognition (ICPR 2012).
- [2] Yuval Netzer, Tao Wang, Adam Coates, Alessandro Bissacco, Bo Wu, and Andrew Y. Ng. “Reading digits in natural images with unsupervised feature learning”. In NIPS Workshop on Deep Learning and Unsupervised Feature Learning, 2011.
- [3] Ian Goodfellow, Yaroslav Bulatov, Julian Ibarz, Sacha Arnoud, Vinay Shet. "Multi-digit Number Recognition from Street View Imagery using Deep Convolutional Neural Networks". ICLR, 2014.
- [4] Fabian Stark , Rudolph Triebel , Daniel Cremers. “CAPTCHA Recognition with Active Deep Learning”. In GCPR Workshop on New Challenges in Neural Computation, 2015.
- [5] Hui Li, Chunhua Shen. “Reading Car License Plates Using Deep Convolutional Neural Networks and LSTMs”. ArXiv technical report, 2016.
- [6] John Duchi, Elad Hazan, Yoram Singer. “Adaptive Subgradient Methods for Online Learning and Stochastic Optimization”. Journal of Machine Learning Research 12, 2011.
- [7] Sebastian Ruder. “An overview of gradient descent optimization algorithms”. <http://sebastianruder.com/>. 2016.
- [8] Karen Simonyan, Andrew Zisserman. “Very Deep Convolutional Networks for Large-Scale Image Recognition”. ArXiv technical report, 2014.